

# Le Protocole MQTT

Nicolas Vadkerti

2 décembre 2019

[https://github.com/SlaynPool/CR\\_LXC-LXD/](https://github.com/SlaynPool/CR_LXC-LXD/)

## 1 Container LXC sosu Linux

### 1.1 Installation de LXC

Pour installer LXC, on doit suivre le TP, càd :

```
apt-get purge lxd
apt-get install lxc lxc-templates lxc-utils bridge-utils debootstrap yum libvirt0
libpam-cgfs
```

Listing 1 – Installation de LXC

Pour vérifier que l'installation fonctionne coorectement on peut utiliser les commandes lxc-...

```
root@ubuntu:~# lxc-checkconfig
Kernel configuration not found at /proc/config.gz; searching...
Kernel configuration found at /boot/config-4.15.0-29-generic
--- Namespaces ---
5 Namespaces: enabled
  Utsname namespace: enabled
  Ipc namespace: enabled
  Pid namespace: enabled
  User namespace: enabled
10 Network namespace: enabled

--- Control groups ---
Cgroups: enabled

15 Cgroup v1 mount points:
  /sys/fs/cgroup/systemd
  /sys/fs/cgroup/rdma
  /sys/fs/cgroup/memory
  /sys/fs/cgroup/cpuset
20 /sys/fs/cgroup/net_cls,net_prio
  /sys/fs/cgroup/devices
  /sys/fs/cgroup/cpu,cpuacct
  /sys/fs/cgroup/hugetlb
  /sys/fs/cgroup/blkio
25 /sys/fs/cgroup/pids
  /sys/fs/cgroup/freezer
  /sys/fs/cgroup/perf_event

Cgroup v2 mount points:
30 /sys/fs/cgroup/unified

Cgroup v1 clone_children flag: enabled
Cgroup device: enabled
Cgroup sched: enabled
35 Cgroup cpu account: enabled
Cgroup memory controller: enabled
Cgroup cpuset: enabled

--- Misc ---
40 Veth pair device: enabled, not loaded
Macvlan: enabled, not loaded
Vlan: enabled, not loaded
Bridges: enabled, loaded
Advanced netfilter: enabled, not loaded
45 CONFIG_NF_NAT_IPV4: enabled, loaded
```

```

CONFIG_NF_NAT_IPV6: enabled, not loaded
CONFIG_IP_NF_TARGET_MASQUERADE: enabled, loaded
CONFIG_IP6_NF_TARGET_MASQUERADE: enabled, not loaded
CONFIG_NETFILTER_XT_TARGET_CHECKSUM: enabled, loaded
50 CONFIG_NETFILTER_XT_MATCH_COMMENT: enabled, not loaded
FUSE (for use with lxcfs): enabled, not loaded

--- Checkpoint/Restore ---
checkpoint restore: enabled
55 CONFIG_FHANDLE: enabled
CONFIG_EVENTFD: enabled
CONFIG_EPOLL: enabled
CONFIG_UNIX_DIAG: enabled
CONFIG_INET_DIAG: enabled
60 CONFIG_PACKET_DIAG: enabled
CONFIG_NETLINK_DIAG: enabled
File capabilities:

Note : Before booting a new kernel, you can check its configuration
65 usage : CONFIG=/path/to/config /usr/bin/lxc-checkconfig

```

Listing 2 – Verification de LXC

Il ne semble pas avoir de soucis, ce qui est plutôt bon signe pour la suite du TP

Si l'on regarde dans `/usr/share/lxc/templates/` on peut voir toutes les distributions que l'on va pouvoir contenariser :

```

root@ubuntu:~# ls /usr/share/lxc/templates/
lxc-alpine      lxc-fedora-legacy  lxc-sabayon
lxc-altlinux    lxc-gentoo         lxc-slackware
lxc-archlinux   lxc-local          lxc-sparclinux
5 lxc-busybox    lxc-oci            lxc-ssh
lxc-centos      lxc-openmandriva   lxc-ubuntu
lxc-cirros      lxc-opensuse       lxc-ubuntu-cloud
lxc-debian      lxc-oracle         lxc-voidlinux
lxc-download    lxc-plamo
10 lxc-fedora     lxc-pld

```

Listing 3 – listes des distributions contenarisables

## 1.2 Création d'un container LXC Debian stretch.

Pour cela, il faut :

```

root@ubuntu:~# lxc-create -t debian -n debian-j1 -- -r stretch -a amd64
ebootstrap is /usr/sbin/debootstrap
Checking cache download in /var/cache/lxc/debian/rootfs-stretch-amd64 ...
gpg: key 7638D0442B90D010: 4 signatures not checked due to missing keys
5 gpg: key 7638D0442B90D010: "Debian Archive Automatic Signing Key (8/jessie) <
  ftpmaster@debian.org>" not changed
gpg: Total number processed: 1
gpg: unchanged: 1
Downloading debian minimal ...
I: Retrieving InRelease
10 I: Retrieving Release
...

root@ubuntu:~# lxc-info debian-j1
Name:      debian-j1
15 State:    STOPPED
root@ubuntu:~# lxc-start debian-j1
root@ubuntu:~# lxc-info debian-j1
Name:      debian-j1
State:     RUNNING
20 PID:      19922
CPU use:   0.12 seconds
BlkIO use: 132.00 KiB
Memory use: 13.07 MiB
KMem use:  2.03 MiB
25 Link:     veth19GBKG
TX bytes:  858 bytes
RX bytes:  780 bytes

```

Listing 4 – premier container

Comme on peut le voir, nous venons de lancer notre premier container via lxc.

Si l'on recree un nouveau container avec un nom different, la vitesse est bien plus rapide. En effet, il ne retelecharge pas toutes l'image, mais copie simplement la précédente. Le cache est ici : `var/cache/lxc/debian/rootfs-stretch-amd64`

Pour déployer d'autre distributions comme centos on peut aussi utiliser LXC

```
lxc-create -t centos -n centos
```

Listing 5 – LXC centos

```
# Lister les containers
$ lxc-ls
centos      debian-j1 debian-j2

5 # Demarrer un containers
$ lxc-start debian-j1

# Stopper Un containers
lxc-stop debian-j1

10 # Redemarrer en mode Deamon

lxc-start -d debian-J1

15 # S'attacher a ce container

root@ubuntu:~# lxc-attach debian-j1
root@debian-j1:~# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen
    1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
20 9: eth0@if10: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group
    default qlen 1000
    link/ether 00:16:3e:aa:9b:c3 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 10.0.3.212/24 brd 10.0.3.255 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::216:3eff:feaa:9bc3/64 scope link
        valid_lft forever preferred_lft forever
30 root@debian-j1:~#

# Voir le PID du Container et les processus fils

35 root@ubuntu:~# lxc-info debian-j1
Name:      debian-j1
State:     RUNNING
PID:       6196
IP:        10.0.3.212
CPU use:   0.25 seconds
BlkIO use: 2.02 MiB
Memory use: 15.79 MiB
KMem use:  2.16 MiB
45 Link:     veth23209A
TX bytes:  1.55 KiB
RX bytes:  1.86 KiB
Total bytes: 3.42 KiB

50 root@ubuntu:~# ps -ejH 6196

6192 6192 6192 ?        Ss      0:00    [lxc monitor] /var/lib/lxc debian-j1
6196 6196 6196 ?        Ss      0:00    /sbin/init
6252 6252 6252 ?        Ss      0:00    /lib/systemd/systemd-journald
55 6295 6295 6295 ?        Ss      0:00    /sbin/dhclient -4 -v -pf /run/dhclient.eth0
    .pid -lf /var/lib/dhcp/dhclient.eth0.leases -I -df /var/lib/dhcp/dhclient6.eth0.
    leases eth0
6331 6331 6331 pts/3    Ss+     0:00    /sbin/agetty --noclear --keep-baud console
    115200,38400,9600 vt220
6332 6332 6332 ?        Ss      0:00    /usr/sbin/sshd -D
```

Listing 6 – Les commandes utiles

```

# Pour limiter l'utilisation Memoire d'un Container
# Dans /var/lib/lxc/$CONTAINER/config
lxc.cgroup.memory.limit_in_bytes = 512M

5 # Pour limiter le Swap du container
memory.memsw.limit_in_bytes= 1G

# Pour verifier, on peut utiliser la commande stress dans le container, et verifier que
  l'on ne peut allouer plus de 512M
stress --vm 1 --vm-bytes 600M

10

# Installer apache2
root@ubuntu:~# lxc-attach debian-j1
root@debian-j1:~# apt install apache2
15 Lecture des listes de paquets... Fait
...
# Faire un autostart
vim /var/lib/lxc/$CONTAINER/config
lxc.start.auto = 1

20

# Figer, relancez le container
root@ubuntu:~# lxc-freeze debian-j1
root@ubuntu:~# lxc-info debian-j1
25 Name:          debian-j1
State:          FROZEN
PID:           1105
IP:            10.0.3.212
CPU use:       0.77 seconds
30 BlkIO use:    21.57 MiB
Memory use:    48.55 MiB
KMem use:      4.82 MiB
Link:          vethG27FGY
TX bytes:      2.10 KiB
35 RX bytes:    16.55 KiB
Total bytes:   18.64 KiB
root@ubuntu:~# lxc-unfreeze debian-j1
root@ubuntu:~# lxc-info debian-j1
40 Name:          debian-j1
State:          RUNNING
PID:           1105
IP:            10.0.3.212
CPU use:       0.77 seconds
BlkIO use:    21.57 MiB
45 Memory use:  48.55 MiB
KMem use:      4.82 MiB
Link:          vethG27FGY
TX bytes:      2.10 KiB
RX bytes:      16.88 KiB
50 Total bytes:  18.98 KiB

# Cloner un container
Il faut que la machine soit stopper puis

55 $: lxc-copy -n debian-j1 -N clonedebian

```

Listing 7 – Les commandes utiles

## 1.3 Comment le container est-il connecté au réseau ?

Le container est relié au réseau grâce à un Bridge Virtuel. On se rend compte grâce à la commande `ip a` que `lxc` monte des interfaces réseaux virtuels. Le fichier relatif aux réseaux est `/etc/default/lxc-net` :




FIGURE 1 – Network Schema

```
root@ubuntu:~# cat /etc/default/lxc-net
# This file is auto-generated by lxc.postinst if it does not
# exist. Customizations will not be overridden.
# Leave USE_LXC_BRIDGE as "true" if you want to use lxcbr0 for your
# containers. Set to "false" if you'll use virbr0 or another existing
# bridge, or mavlan to your host's NIC.
USE_LXC_BRIDGE="true"

# If you change the LXC_BRIDGE to something other than lxcbr0, then
# you will also need to update your /etc/lxc/default.conf as well as the
# configuration (/var/lib/lxc/<container>/config) for any containers
# already created using the default config to reflect the new bridge
# name.
# If you have the dnsmasq daemon installed, you'll also have to update
# /etc/dnsmasq.d/lxc and restart the system wide dnsmasq daemon.
LXC_BRIDGE="lxcbr0"
LXC_ADDR="10.0.3.1"
LXC_NETMASK="255.255.255.0"
LXC_NETWORK="10.0.3.0/24"
LXC_DHCP_RANGE="10.0.3.2,10.0.3.254"
LXC_DHCP_MAX="253"
# Uncomment the next line if you'd like to use a conf-file for the lxcbr0
# dnsmasq. For instance, you can use 'dhcp-host=maill,10.0.3.100' to have
# container 'maill' always get ip address 10.0.3.100.
#LXC_DHCP_CONFILE=/etc/lxc/dnsmasq.conf

# Uncomment the next line if you want lxcbr0's dnsmasq to resolve the .lxc
# domain. You can then add "server=/lxc/10.0.3.1" (or your actual $LXC_ADDR)
# to your system dnsmasq configuration file (normally /etc/dnsmasq.conf,
# or /etc/NetworkManager/dnsmasq.d/lxc.conf on systems that use NetworkManager).
# Once these changes are made, restart the lxc-net and network-manager services.
# 'container1.lxc' will then resolve on your host.
#LXC_DOMAIN="lxc"
```

Listing 8 – lxc-net default

Vu du Container, cela ressemble a ceci :

```
9: eth0@if10: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group
  default qlen 1000
    link/ether 00:16:3e:aa:9b:c3 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 10.0.3.212/24 brd 10.0.3.255 scope global eth0
      valid_lft forever preferred_lft forever
    inet6 fe80::216:3eff:feaa:9bc3/64 scope link
      valid_lft forever preferred_lft forever
```

Listing 9 – Configuration Réseau

## 2 Création des CGROUPS

### 2.1 Les xterm consommateurs

```

top - 17:19:34 up 19 min,  2 users,  load average: 1,42, 0,58, 0,32
Tasks:  85 total,   3 running,  82 sleeping,   0 stopped,   0 zombie
%Cpu(s): 23,3 us, 76,7 sy,  0,0 ni,  0,0 id,  0,0 wa,  0,0 hi,  0,0 si,  0,0 st
MiB Mem :   987,3 total,   156,4 free,    97,4 used,   733,6 buff/cache
5 MiB Swap:   488,0 total,   481,0 free,    7,0 used.   723,7 avail Mem

   PID USER      PR  NI    VIRT    RES    SHR S  %CPU  %MEM     TIME+ COMMAND
10  25488 user       20   0    5272    748    688 R   49,8   0,1   0:50.58 md5sum
   25490 user       20   0    5272    752    688 R   49,8   0,1   0:28.06 md5sum
   24206 user       20   0   17024   5124   4028 S    0,3   0,5   0:00.58 sshd

```

Listing 10 – top

On peut voir que les deux processus prennent 50 % du CPU de la vm environ.

Pour limiter l'utilisation de ou augmenter l'utilisation du CPU, on peut utiliser les cgroupes comme ceci

```

#!/bin/bash

cgcreate -g cpu,cpuset:/quatrevingtpourcentcpu
5 cgcreate -g cpu,cpuset:/vingtpourcent
  cpucgset -r cpu.shares=2 vingtpourcent
  cpucgset -r cpu.shares=8 quatrevingtpourcent
  cpucgget -r cpu.shares quatrevingtpourcent
  cpucgget -r cpu.shares vingtpourcent
10 cpucgexec -g cpu:/quatrevingtpourcentcpu
  xterm -bg orange -e "md5sum /dev/urandom" &
  cgexec -g cpu:/vingtpourcentcpu xterm -bg blue -e "md5sum /dev/urandom" &
  top -d2

```

Listing 11 – Cgroupe set

Pour utiliser cgexec, il faut être root, et comme root ne peut utiliser de session graphique en temps normal, il faut modifier etc/pam.d/su et /etc/pam.d/su-l en rajoutant “latex veut pas” pour que ça fonctionne.

On

```

top - 17:35:13 up 34 min,  2 users,  load average: 0,56, 0,41, 0,56
Tasks:  86 total,   4 running,  82 sleeping,   0 stopped,   0 zombie
%Cpu(s): 28,3 us, 70,9 sy,  0,0 ni,  0,0 id,  0,0 wa,  0,0 hi,  0,8 si,  0,0 st
MiB Mem :   987,3 total,   144,6 free,    99,5 used,   743,3 buff/cache
5 MiB Swap:   488,0 total,   481,2 free,    6,8 used.   721,4 avail Mem

   PID USER      PR  NI    VIRT    RES    SHR S  %CPU  %MEM     TIME+ COMMAND
10  25672 root       20   0    5272    752    688 R   79,8   0,1   0:31.67 md5sum
   25676 root       20   0    5272    688    624 R   19,9   0,1   0:02.67 md5sum
     1 root       20   0   105164  10112   7500 S    0,0   1,0   0:03.00 systemd

```

Listing 12 – top